

Priority map

Priority	Type	What to change
Critical	Code + paper	Add stronger evidence against the Linear Regression baseline: report the number of test points satisfying $y \geq 0.20$ $y \geq 0.20$, run multiple seeds or repeated evaluation, and include a significance test for LR vs proposed; the paper currently reports 1,266 test sequences after a 48-hour lookback and gives peak RMSE in Table 7, but no uncertainty or sample-count support for the peak claim.
Critical	Code + paper	Rework the scheduling section so the operational gain is attributed to forecast quality, not just the threshold rule; Table 11 currently reports a 0.15 threshold, 280 eligible hours out of 1,266, and renewable-share improvement from 15.34 to 25.81 without comparing LR vs proposed vs perfect-foresight under the same rule.
Critical	Code + paper	Either provide quantitative blockchain prototype evaluation or explicitly downgrade the blockchain part to a proposed architecture; the manuscript describes IPFS/CID anchoring, keccak256 hashing, smart contracts, and a 7,000/10,000 PGS gate, but it also says gas cost, latency, privacy, and consensus evaluation remain future work.
Critical	Paper, or remove section	Define the Predictive Green Score with an exact formula and calibration procedure; the introduction describes a 0–100 confidence score, while the system section uses an on-chain

		7,000/10,000 threshold, and the future-work section still says calibrated confidence scores need to be incorporated.
Critical	Paper, unless justified by code	Clarify or delete the “optimized oracle export” in Table 8; the paper reports 3,360 rows and RMSE 0.000170 for the optimized oracle export, which is dramatically smaller than the main test RMSE 0.007337, but the mechanism is not clearly defined in the manuscript.
High	Paper	Rewrite abstract, results, discussion, and conclusion to state parity with Linear Regression on global metrics; Table 6 shows LR test RMSE/R ² of 0.007334/0.988058 and the proposed model at 0.007337/0.988050, while the abstract still highlights 25.3% better RMSE and 44.2% lower unexplained variance versus Random Forest.
High	Code + paper	Add stronger contemporary baselines if feasible; the current benchmark set includes LSTM, GRU, BiLSTM, XGBoost, Linear Regression, and Random Forest, while the literature review discusses transformer-family models without benchmarking them.
High	Paper	Unify naming everywhere; the manuscript alternates among labels such as “residual decomposition-based model,” “proposed residual,” “attention-based BiLSTM residual learner,” and historical EA-BiLSTM variants.
High	Paper	Reconcile internal inconsistencies across tables and text; examples include best validation loss shown as 4.23×10^{-5} in Table 4 versus 4.2252×10^{-5} in the text, and standalone BiLSTM results differing from EA-BiLSTM v1 results across Tables 5, 6, 9, and 10.

Medium	Paper	Fix the implementation description in the introduction; the intro says the RNN works on “minutes to hours” and “10–15 minutes,” but the methodology clearly states a 48-hour lookback sequence with 10 input channels.
Medium	Paper	Remove or rephrase the near-zero residual-mean argument; the manuscript currently presents the training residual mean “approximately 0.000000” as evidence, which is not a meaningful empirical result in this setup.
Medium	Paper	Tighten writing and presentation; the manuscript contains repeated explanations across results/discussion, the typo “REASEARCH GAPS,” and inconsistent capitalization such as LSTM/BiLSTM variants and mid-sentence all-caps terms.